

Evaluation frame



PDF

294 pages

229 government + 65 arXiv



Held-out Q-A probe

663 Q-A

527 government + 136 arXiv

query + reference page
+ answer span

Candidate pipelines

Parser candidates P_i

P_1 P_2 ... P_m

Parsed page outputs

Intrinsic
score

Chunker candidates C_j

C_1 C_2 ... C_c

Intrinsic
ranking

all $P_i \times C_j$ pairs

Compare
Rankings
(same chunker)

C1

Intrinsic
ranking
can mislead

Candidate corpora

$C_j(P_i)$, $m \times c$

Index each corpus

candidate index

RCPS evaluation

Search with probe queries

BGE-M3

mE5-large

Qwen3-Emb-8B

Relevant chunk

from reference page AND contains
normalised answer span

MRR@k

RCPS(P, C)

mean MRR across retrievers and cutoffs

C2

Controlled settings

Probe set

Retriever models

Cutoffs $k = 1, 5, 10$

Relevance rule

Candidate-wise RCPS scores

one score for every $P_i \times C_j$

Select, diagnose + act

RCPS score matrix

	C_1	...	C_j	...	C_c
P_1					
...					
P_i			✓		
...					
P_m					

RCPS rank | latency/compute tie-break

Selected $P^* + C^*$

parser + chunker configuration

P^* output

C^* chunks

Ref. page
+ span

Coverage diagnostic

retriever-free exact-span check

C3

Absent
no match
in output

Covered
one chunk

Split
across chunks

Inspect /
switch P

Proceed

Overlap /
rechunk

Parser training
(optional)

Re-evaluate
changed P/C

Deploy final $P^* + C^*$