

Parser / chunker
candidates

Held-out Q-A probe
(~few hundred, no training)

chunk the parsed corpus

the protocol (not a new metric)

(iii) format-invariant relevance:

chunk relevant iff it contains the gold answer span

(i) extrinsic: score retrieval on the probe, not the text

(ii) retriever-agnostic: average $\text{MRR}@ \{1, 5, 10\}$
over $R = \{\text{BGE-M3, mE5-large, Qwen3-Emb-8B}\}$

RCPS(P) = retriever-averaged MRR

rank & select parser / chunker
(picks what TEDS / BC rank wrong)

0.20 → 0.55 Hit@1 on our data