

Parser / chunker
candidates

Held-out Q-A probe
(few hundred, no training)

chunk the parsed corpus

(i) extrinsic — score retrieval on the probe, not the text

fixes: text scores (edit dist., BC) mispredict retrieval

(ii) retriever-averaged MRR over {BGE-M3, mE5, Qwen3-Emb}

fixes: a single embedder can flip the near-tied top parser

(iii) format-normalised relevance — chunk holds the answer span

fixes: credits content the parser kept, not its formatting

RCPS(P) = retriever-averaged MRR

rank & select parser / chunker
(picks what edit dist. / BC rank wrong)

Hit@1 0.20 → 0.55 on our data